



BF16 training, same data and inherited Adam; branch-specific total objectives omitted

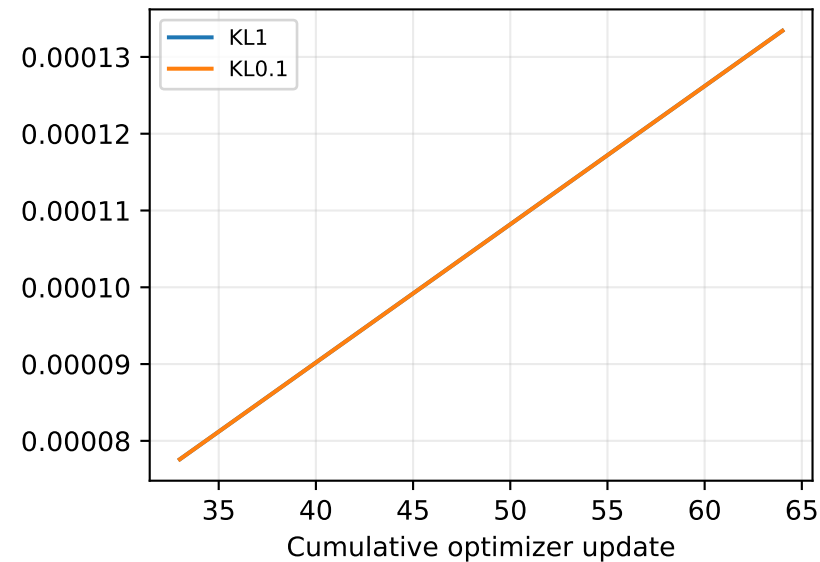
Pre-clip gradient norm



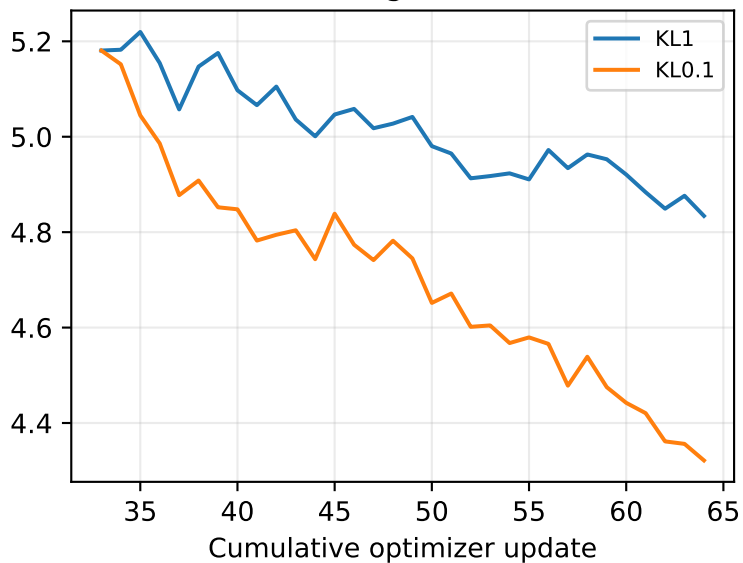
Clipping coefficient estimate



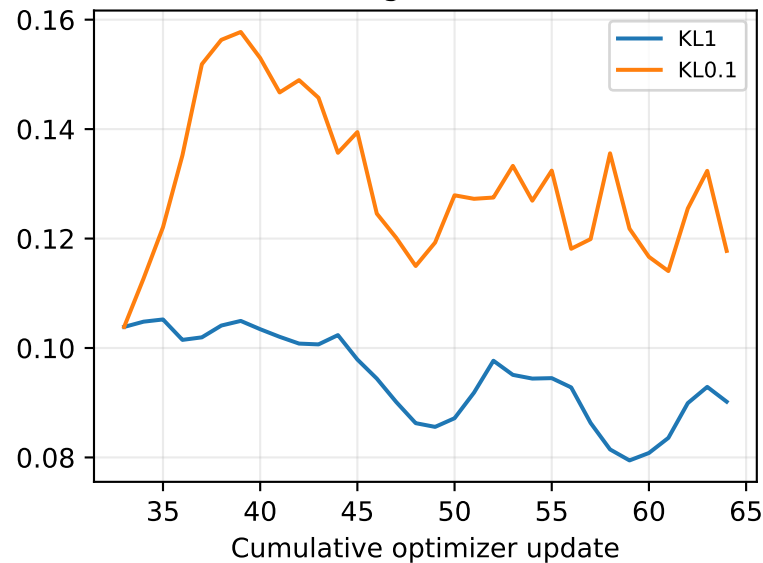
Learning rate used



Training raw CE



Training raw latent



Training raw KL

