



Pretrained OLMo-1B: complete optimizer-step throughput

B16 / T2048



B64 / T512



H100 80GB · BF16 mixed · CUDA graphs · full CE · five timed updates/run

Whiskers: range of run medians, not confidence intervals. See retained numerical qualifications; no quality claim.