



Physical batch per GPU

B64 B96

Full-update throughput



Peak allocated memory, including setup



T512; three-update medians. RT uses layers (0,15); FBT uses K2.

\* RT + FBT retains its failed initial gradient screen; resource results do not clear it.