



Loss ramp uses phase updates; Fuzzy CE is dense, A5 CE scores every state

Task objectives are averaged separately before weighting; NextLat weight remains one


Fuzzy loss coefficient



Weighted total training objective



Unweighted task CE



Unweighted task NextLat SmoothL1

