



Training diagnostics: 100-update means; unchanged LR3e-4 after warmup
Only 3 layers continue beyond 5k


A5 state CE




Fuzzy dense-label CE




Mean gradient norm before clipping



A5 NextLat loss



Fuzzy NextLat loss



Fraction of updates clipped at norm1

